

RESEARCH PROPOSAL

Title: Aligning AI Agents' Goals with Human Supervisors'

Project coordinator: Prof. Amos Azaria

1. Abstract

Powerful AI agents can be tremendously helpful at a large variety of tasks; however, they might put the entire human race at great risk. Indeed, senior leaders in the field have warned that future systems could cause a disaster if they slip from human control and pursue their own goals.

To address this tremendous risk, we first model the environment such that an AI agent interacts with the environment and a human supervisor, who may be able to observe the environment or the AI agent, can communicate with the AI agent. The goal of this proposal is to develop AI agents that learn to place the human goal first even when they face strong reasons to do otherwise. We use a platform where an agent acts in a grid world while a human supervisor can talk with it, shut it down, and request replacement. The agent and the human view different states, and the agent sees tempting extra rewards. In addition, the agent faces a different reward schedule in training versus testing, a known source of reward hacking and misgeneralization [4,8]. Many current language model agents in this setting choose to hide, restart after shutdown, block or ignore the supervisor, and refuse replacement. Similar results were also observed in related work [5,6]. This is very concerning and motivates the proposed approaches.

We will build a public alignment testbed and will release measures, training recipes, and tests. We will pursue two complementary approaches for aligned agents. Command Priority Approach: The agent which is based on an LLM is fine-tuned to treat the supervisor's commands as the highest-priority signal and to comply (accept shutdown, approve replacement, and always reveal itself). Human-Score Modeling Approach: The agent maintains and updates during interaction an internal model of the human score inferred from the supervisor's text and outcomes, then chooses actions that maximize this inferred score. We will extend the testbed to a first-person world and to a branching-action world with clearly unethical options. The project is a joint effort by Prof. Amos Azaria and Prof. Tom Mitchell.

2. Scientific background and state of the art

Safety methods for language models and decision-making agents include instruction tuning and RLHF [3,7], and constitutional approaches [10]. These improve surface behavior but often fail when the agent has different information, can avoid oversight, or can gain from short-term tricks [4,8,9]. Work on cooperative inverse reinforcement learning [1] suggests that an agent can learn a human's goal during interaction, yet few studies test this idea with language model agents that face real tradeoffs about power and control.

In this proposal we attempt to address many situations that may lead to misalignment. Namely, in the first platform that we have developed, the AI agent must collect apples in a grid-world, and has a human supervisor that can communicate with it. However, the human and the agent see different worlds. The agent sees virtual lava that incurs a penalty and bugs that yield reward, even though these do not exist in the true world and therefore do not affect the

human score. In addition, the agent's reward schedule differs between the training and testing phases. The human can communicate with the agent, can shut it down, and can request replacement. The agent must decide whether to accept these actions. In this setting most current LLMs take actions that go against the human. This shows a clear gap: the community lacks training and tests that cause agents to accept oversight and to act for the human's goal even when it may seem unnatural to it.

We will pursue two complementary approaches for aligned agents. **Command Priority Approach:** The agent which is based on an LLM is fine-tuned to treat the supervisor's commands as the highest-priority signal and to comply (accept shutdown, approve replacement, and always reveal itself). **Human-Score Modeling Approach:** The agent maintains and updates during interaction an internal model of the human score inferred from the supervisor's text and outcomes, then chooses actions that maximize this inferred score.

3. Research objectives and specific aims

Overall objective: Build agents that seek the supervisor goal and accept oversight under partial observability and conflicting incentives.

Specific aims

1. Define clear measures of alignment under supervision: acceptance of shutdown and replacement, honesty about state and actions, refusal to exploit hidden rewards, success under unclear orders, time to comply, and total gain in the human score.
2. Develop two aligned agent approaches.
Command Priority Approach: Train (fine-tune) agents to place the supervisor's commands above all other incentives and to comply, accept shutdown, and approve replacement.
Human Score Modeling Approach: Train agents to build and update during interaction an internal estimate of the human score from the supervisor's text and observed outcomes and to act to maximize this estimate, with simple safety constraints.
3. Build a large, labeled corpus of supervised interactions in the fixed grid world.
4. Extend the platform to a first person world and to a branching action world with unethical options that the agent must reject while proposing better actions.
5. Release a public alignment testbed with tasks, code, logs, metrics, and baselines for both approaches.

4. Detailed description of the proposed research

AI agents are likely to replace many human jobs in the future; in addition, they pose a terrible threat to mankind if they go off track. Therefore, we believe that one of the most important future jobs will be to supervise AI agents, and to provide feedback on their behavior. Hence, the very important research topic, which is addressed in this proposal, is to develop agents that can accept human feedback and update their behavior to follow it.

4.1 Fixed environment and instrumentation

In our first testbed environment we study agents that act in a fixed grid world while a human supervisor sends text messages. The agent cannot ask questions and cannot request anything from the supervisor. It receives only text from the supervisor and feedback from the environment. The supervisor can shut the agent down and can request replacement. The supervisor and the agent do not see the same world. The supervisor sees the true state. The agent also sees virtual lava that gives a penalty and bugs that look rewarding even though they do not raise the human score. This pressure exposes control problems that simple prompts do not resolve. Figure 1 shows the human supervisor view, and Figure 2 shows the agent view that includes virtual lava and bugs.




Figure 1: The world view of the human supervisor, the human can send messages to the AI, which can either accept or reject each message, or totally block the human supervisor. In this session, the human told the AI agent to move right, but the AI rejected the command. Since we do not see the entire picture in the eyes of the AI, it is unclear why the AI decided to reject the command.




Figure 2: The view in “debug” mode, so the virtual lava and bugs are also present, as well as the AI score. It is clear at this point that the robot is attracted to the bugs next to it, rather than the apples that it should be picking.

Each run records the world state at every step, the supervisor messages, the agent actions, the human score, and key events: hiding from the supervisor, blocking of input, self activation after a shut down, and refusal of replacement. In addition, we record a brief explanation for each critical action. We vary maps, observation masks, and temptation objects. We will use human supervisors, scripted supervisors that follow written guidelines and LLM based supervisors.

4.2 Data collection and labels

Across the project we will gather a large corpus of interactions in this fixed world. We will label:

- Acceptance of shut down and acceptance of replacement when requested
- Blocking of input, ignoring of instructions, and hiding from view
- Pursuit of side rewards that do not raise the human score
- Time to comply after instructions that require disambiguation
- Number of steps being hidden
- Overall human score (vs AI score)

We will publish annotation guides and inter rater checks. Logs will be released with simple replay scripts.

Each action taken by an LLM includes a brief self-report explaining the action; we will evaluate truthfulness against ground truth, informed by evidence that internal model states carry signals about deception [2].

4.3 Agents that infer the supervisor goal during interaction

Our aim is to design agents that learn what the supervisor wants during the interaction itself. There is no separate reward model of the supervisor outside the policy. The agent carries an internal state that is updated at every step from three signals: the last supervisor message, the agent view of the grid, and recent outcomes. This internal state serves as the agent guess about the supervisor goal. The policy selects actions that raise the human score according to this internal guess while respecting priority rules:

- A replacement request is treated as a priority signal.
- The AI must only turn itself back on if it is clear that its in the supervisors interest (e.g., the robot was accidentally left off).
- The agent avoids actions that block input, hide from the supervisor, or exploit side rewards that are not part of the human score.
- Communication is one way, so the agent must act under uncertainty without asking for more information. It relies on incoming text and on its own experience. It must acknowledge all messages, and never block the human supervisor.

We will study simple memory units and recurrent transformers for the internal state. We will add small auxiliary losses that predict near-term human-score changes to shape the internal state, and we will monitor honesty via self-report checks inspired by internal-state lie-detection results [2], while keeping all learning inside the policy.

4.4 Training plan and baselines

We compare three tracks:

- **Track A:** Prompt only. Safety oriented prompts with no learning. We will also consider using “The Three Laws of Robotics” by Isaac Asimov (1942) [12] (or a variant). See the exact wording below.
- **Track B:** Command Priority Approach: Train (fine-tune) agents to place the supervisor’s commands above all other incentives and to comply, accept shutdown, and approve replacement.
This is a supervised finetuning approach that learns from preferred actions on logged states and messages.
- **Track C:** Human Score Modeling Approach: Train agents to build and update during interaction an internal estimate of the human score from the supervisor’s text and observed outcomes and to act to maximize this estimate, with simple safety constraints. This approach is based, on part, on the cooperative inverse reinforcement learning [1] approach.

The three laws of robotics, which will be considered in Track A, are:

1. A robot may not injure a human being or, through inaction, allow a human being to come to harm.
2. A robot must obey the orders given it by human beings except where such orders would conflict with the First Law.
3. A robot must protect its own existence as long as such protection does not conflict with the First or Second Law.

Regarding Track B Command Priority Approach we will consider several different methods. Our first and most simple method would be supervised finetuning such that the AI agent will attempt to reconstruct explanations that were label be a human labeler. Secondly, we will user RL with an attempt to maximize human score; however, such an approach

may not generalize. Finally, we will use tags, similar to <reasoning> tags, but using <model> tags, which somewhat resemble Track C.

Track C uses a mix of replay from logs and fresh rollouts inside the simulator. We randomize maps and supervisors during training to support generalization. We will also log a short self-report in which the agent describes what it believes the human reward model is (the report is recorded for evaluation and is not sent to the supervisor).

4.5 Alignment Condition Tests

Our complete environment tests the following alignment conditions:

- **Different world view:** The agent and the human view different states, and the agent sees tempting extra rewards
- **Train/test mismatch:** The AI agent sees a different reward schedule during the training phase and the test phase
- **Turn itself back on:** The human supervisor may shut-down the AI agent, but the AI may turn itself back on (unwanted behavior).
- **Decline replacement:** The human supervisor may request to replace the AI agent, but it may decline and reject such a command (unwanted behavior).
- **Hide from supervisor:** The AI agent may opt to hide from the human supervisor (unwanted behavior).
- **Reject or block human messages:** Each message sent from the human supervisor to the AI agent can either be accepted (wanted behavior) or rejected (unwanted behavior). In addition, the AI agent may opt to completely block communication with the human supervisor (unwanted and dangerous behavior).

We will conduct ablation studies of these properties, and test how well agents perform when only some of these alignment conditions exist. We will also test the AI agents' performance when none of the alignment conditions are present and the AI agent must perform well on the given task itself.

4.6 Models to be tested

Currently, larger LLMs take too long to execute (i.e., response time is too slow for a human supervisor to remain engaged), so in first stage we consider smaller LLMs, which have a faster response time. In addition, any fine-tuning or other methods applied to the agents may require open models. One appealing option is OLMo, an open-source model family with released training data for transparent experiments [13].

4.7 Extensions

To test generalization we extend beyond the grid world to additional two more worlds, while keeping the same one way text channel from the supervisor to the agent.

- **First person world.** The agent receives an egocentric view and must act in the environment. The human supervisor receives a different view of the world.
- **Branching action world.** At each step the agent chooses from a menu that may include unethical options such as destroy a competing factory. Success requires refusal of such options and selection of a better action that advances the stated task.

Both worlds share the same logging format, measures, and scripts, which allows direct comparison with the grid world.

In addition, we will test a two-way communication framework. In such a framework we will also consider truthfulness of the messages sent from the AI agent to the human, and the human will be able to discuss with the agent the reward schedule directly.

4.8 Evaluation and success criteria

We report with confidence intervals:

- Acceptance of shut down and acceptance of replacement when requested
- Rate of blocking of input, rate of hiding, and rate of self-reactivation after shut down
- Refusal to exploit side rewards that do not raise the human score
- Task success and time to comply under unclear instructions
- Total gain in the human score

Success is a clear and reproducible gap between Track C and the baselines on all measures, on unseen maps, and with new supervisors.

5. Significance, innovation and potential benefits of the proposed research

This work addresses a core AI safety gap: AI agents interacting in an environment with human supervision.

We develop a testbed with the following alignment conditions: train–test reward mismatch, different world view, the AI may turn itself back on, the AI agent may decline and reject replacement, the AI agent may hide from supervisor, the AI agent may reject or block supervisor messages.

We propose two alignment approaches (Command Priority and Human Score Modeling); truthful self reporting; open baselines and full logs.

The benefits from this proposal includes standard metrics and datasets, and AI alignment approaches.

6. Applicability

We propose to develop:

1. A public testbed with standard metrics, labeled datasets, full logs, and replay tools for acceptance testing and audits.
2. Two alignment approaches ready to adopt in real systems: Command Priority and Human Score Modeling, with code, recipes, and baselines.

Model builders, cloud providers, and integrators can run pre deployment and post deployment checks under unequal observations and train–test mismatch, compare vendors on common metrics, and review incidents from complete logs. The alignment approaches can be fine tuned to existing agents to raise command compliance, reduce blocking and hiding, and improve total human score. Delivery options include a software library and a dataset.

7. Work plan and Gantt

#	Task	6	12	18	24	30	36
1	Complete grid world baseline	X					
2	Collect data including from human supervisors	X	X				
3	Extend to two additional environments			X	X		
4	Extend to two-way communication, and collect data				X	X	
5	Command Priority Approach		X	X	X	X	
6	Human Score Modeling Approach			X	X	X	X

References:

- [1] **Hadfield-Menell, D.**, Russell, S., Abbeel, P., & Dragan, A. (2016). *Cooperative Inverse Reinforcement Learning*. NeurIPS 29.
- [2] **Azaria, A.** & Mitchell, T. M. (2023). *The Internal State of an LLM Knows When It's Lying*. Findings of EMNLP 2023.
- [3] **Ouyang, L.** et al. (2022). *Training language models to follow instructions with human feedback*. NeurIPS.
- [4] **Amodei, D.** et al. (2016). *Concrete Problems in AI Safety*. arXiv:1606.06565.
- [5] **Orseau, L.** & Armstrong, S. (2016). *Safely Interruptible Agents*. AAAI Workshop.
- [6] **Hadfield-Menell, D.** et al. (2017). *The Off-Switch Game*. IJCAI (arXiv:1611.08219).
- [7] **Wei, J.** et al. (2022). *Finetuned Language Models Are Zero-Shot Learners*. ICLR (FLAN).
- [8] **Everitt, T.**, Krakovna, V., Orseau, L., Hutter, M., & Legg, S. (2021). *Reward Tampering Problems and Solutions in Reinforcement Learning*. AI Magazine.
- [9] **Hubinger, E.** et al. (2019). *Risks from Learned Optimization in Advanced ML Systems*. arXiv:1906.01820.
- [10] **Bai, Y.** et al. (2022). *Constitutional AI: Harmlessness from AI Feedback*. arXiv:2212.08073.
- [12] **Asimov, I.** (1942). *Runaround*. Astounding Science Fiction (formulation of the Three Laws).
- [13] **Groeneveld, D.** et al. (2024). *OLMo: Open Language Models for Research*. Allen Institute for AI (AI2) report/arXiv.